

Introduction to Artificial Intelligence

HW3-Report

0716325-曾正豪

Part.1: Minimax Search

In this part, I follow the instruction in the slide of Lec6 p.29.

```
131 # Begin your code (Part 1)
132 act, val = self.value(gameState, 0, 0)
133 return act
```

```
137 def value(self, gameState, player, current_depth):
138     ...
139     return (action, value)
140     ...
141     if current_depth >= self.depth and player == 0:
142         return 0, self.evaluationFunction(gameState)
143     if gameState.isWin() or gameState.isLose():
144         return 0, self.evaluationFunction(gameState)
145
146     if player == 0:
147         current_depth += 1
148         return self.max_value(gameState, player, current_depth)
149     else:
150         return self.min_value(gameState, player, current_depth)
151
152     return 0
```

The first part is function `value()`. Since it need the action for the `getAction` function, and value for `max_value` function and `min_value` function. So I modify it to return the action and value at the same time. It's parameter is `gameState` for the state of current game, `player` for which player currently deciding action, `current_depth` indicate the depth of the search currently. In the `getAction` function, I call it with `player = 0` (pacman) and `depth = 0`.

In line 141 to 144, it will check if it reach the terminal state, and will return the evaluation function value. In the first "if", it check whether it reach the maximum depth. In the second "if" it check if the game is over.

In line 146 to 148 if player turns to pacman, indicate that the depth should plus 1, and return the maximum value of child nodes. Otherwise, it's ghost's turn, it should return the minimum value of child nodes in line 149 to 150.

```

154     def max_value(self, gameState, player, current_depth):
155         ...
156         return (action, value)
157         ...
158         v = -99999999
159         ret_act = None
160         action_list = gameState.getLegalActions(player)
161         for action in action_list:
162             next_state = gameState.getNextState(player, action)
163             next_player = (player + 1) % gameState.getNumAgents()
164             act, val = self.value(next_state, next_player, current_depth)
165             if val > v:
166                 v = val
167                 ret_act = action
168
169         return ret_act, v

```

In the `max_value` function, its parameters' meaning is exactly the same as in the `value` function.

Firstly, I initialize the value to $-\infty$ (I think -99999999 is big enough), and I get all the possible action of the player. For all the action in the action list, I firstly get the next state and next player respectively. And send it into the `value` function to evaluate the value of this child node. Finally, if the new value is bigger than `v`, I will update the new value and action to return.

```

171     def min_value(self, gameState, player, current_depth):
172         ...
173         return (action, value)
174         ...
175         v = 99999999
176         ret_act = None
177         action_list = gameState.getLegalActions(player)
178         for action in action_list:
179             next_state = gameState.getNextState(player, action)
180             next_player = (player + 1) % gameState.getNumAgents()
181             act, val = self.value(next_state, next_player, current_depth)
182             if val < v:
183                 v = val
184                 ret_act = action
185
186         return ret_act, v
187
188     # End your code (Part 1)

```

In the `min_value` function, its parameters' meaning is exactly the same as in the `value` function.

Firstly, I initialize the value to ∞ (I think 99999999 is big enough), and I get all the possible action of the player. For all the action in the action list, I firstly get the next state and next player respectively. And send it into the `value` function to evaluate the value of this child node. Finally, if the new value is smaller than `v`, I will update the new value and action to return.

Part.2: Alpha-Beta Pruning

In this part, I modify the code from part.1 by adding parameters alpha and beta.

```
201         # Begin your code (Part 2)
202         act, val = self.value(gameState, 0, 0, -99999999, 99999999)
203         return act
```

```
206     def value(self, gameState, player, current_depth, alpha, beta):
207         ...
208         return (action, value)
209         ...
210         if current_depth >= self.depth and player == 0:
211             return 0, self.evaluationFunction(gameState)
212         if gameState.isWin() or gameState.isLose():
213             return 0, self.evaluationFunction(gameState)
214
215         if player == 0:
216             current_depth += 1
217             return self.max_value(gameState, player, current_depth, alpha, beta)
218         else:
219             return self.min_value(gameState, player, current_depth, alpha, beta)
220
221         return 0
```

The first part is function value(). I add 2 parameters, alpha and beta. Alpha is for the alpha in Alpha-Beta Pruning, beta is for the beta in Alpha-Beta Pruning. In the getAction function, I call it with player = 0 (pacman) and depth = 0, alpha = -99999999 for $-\infty$, beta = 99999999 for ∞ .

In line 210 to 213, it will check if it reach the terminal state, and will return the evaluation function value. In the first “if”, it check whether it reach the maximum depth. In the second “if” it check if the game is over.

In line 215 to 217 if player turns to pacman, indicate that the depth should plus 1, and return the maximum value of child nodes. Otherwise, it's ghost's turn, it should return the minimum value of child nodes in line 218 to 219.

```
223     def max_value(self, gameState, player, current_depth, alpha, beta):
224         ...
225         return (action, value)
226         ...
227         v = -99999999
228         ret_act = None
229         action_list = gameState.getLegalActions(player)
230         for action in action_list:
231             next_state = gameState.getNextState(player, action)
232             next_player = (player + 1) % gameState.getNumAgents()
233             act, val = self.value(next_state, next_player, current_depth, alpha, beta)
234             if val > v:
235                 v = val
236                 ret_act = action
237             if v > beta:
238                 return ret_act, v
239             alpha = max(alpha, v)
240
241         return ret_act, v
```

In the `max_value` function, its parameters' meaning is exactly the same as in the value function.

Firstly, I initialize the value to $-\infty$ (I think -99999999 is big enough), and I get all the possible action of the player. For all the action in the action list, I firstly get the next state and next player respectively. And send it into the value function to evaluate the value of this child node. then, if the new value is bigger than `v`, I will update the new value and action to return. Next is to check if `v` is bigger than `beta`, if yes, means it need to prune the remaining child nodes, and return immediately. Otherwise after that, it update the `alpha` by `max(alpha, v)`

```
243 def min_value(self, gameState, player, current_depth, alpha, beta):
244     ...
245     return (action, value)
246     ...
247     v = 99999999
248     ret_act = None
249     action_list = gameState.getLegalActions(player)
250     for action in action_list:
251         next_state = gameState.getNextState(player, action)
252         next_player = (player + 1) % gameState.getNumAgents()
253         act, val = self.value(next_state, next_player, current_depth, alpha, beta)
254         if val < v:
255             v = val
256             ret_act = action
257         if v < alpha:
258             return ret_act, v
259         beta = min(beta, v)
260
261     return ret_act, v
262
263 # End your code (Part 2)
```

In the `min_value` function, its parameters' meaning is exactly the same as in the value function.

Firstly, I initialize the value to $-\infty$ (I think -99999999 is big enough), and I get all the possible action of the player. For all the action in the action list, I firstly get the next state and next player respectively. And send it into the value function to evaluate the value of this child node. then, if the new value is smaller than `v`, I will update the new value and action to return. Next is to check if `v` is smaller than `alpha`, if yes, means it need to prune the remaining child nodes, and return immediately. Otherwise after that, it update the `beta` by `min(beta, v)`

Part.3: Expectimax Search

In this part, I modify the code from part.1 by replacing the min_value function by expect_value function.

```
277     # Begin your code (Part 3)
278     act, val = self.value(gameState, 0, 0)
279     return act

281     def value(self, gameState, player, current_depth):
282         ...
283         return (action, value)
284         ...
285         if current_depth >= self.depth and player == 0:
286             return 0, self.evaluationFunction(gameState)
287         if gameState.isWin() or gameState.isLose():
288             return 0, self.evaluationFunction(gameState)
289
290         if player == 0:
291             current_depth += 1
292             return self.max_value(gameState, player, current_depth)
293         else:
294             return self.expect_value(gameState, player, current_depth)
295
296         return 0
```

The first part is function value(). It is almost the same as in part 1, besides the 294 line I replace the min_value function by expect_value function

```
298     def max_value(self, gameState, player, current_depth):
299         ...
300         return (action, value)
301         ...
302         v = -99999999
303         ret_act = None
304         action_list = gameState.getLegalActions(player)
305         for action in action_list:
306             next_state = gameState.getNextState(player, action)
307             next_player = (player + 1) % gameState.getNumAgents()
308             act, val = self.value(next_state, next_player, current_depth)
309             if val > v:
310                 v = val
311                 ret_act = action
312
313         return ret_act, v
```

The max_value function is the same as in part.1

```

315     def expect_value(self, gameState, player, current_depth):
316         ...
317         return (action, value)
318         ...
319         ret_act = None
320         action_list = gameState.getLegalActions(player)
321         total = 0
322         for action in action_list:
323             next_state = gameState.getNextState(player, action)
324             next_player = (player + 1) % gameState.getNumAgents()
325             act, val = self.value(next_state, next_player, current_depth)
326             total += val
327
328         return ret_act, total / len(action_list)
329
330     # End your code (Part 3)

```

In the `expect_value` function, its parameters' meaning is exactly the same as in the `value` function.

Firstly, I initialize the *total* to 0, and I get all the possible action of the player. For all the action in the action list, I firstly get the next state and next player respectively. And send it into the value function to evaluate the value of this child node. Lastly, I add the child node's value to *total*. In the end of this function, I return the *total* divided by the length of possible actions since it is uniformly at random.

Part.4: Evaluation Function (Bonus)

My result:

```

Question part4
Pacman emerges victorious! Score: 1154
Pacman emerges victorious! Score: 1147
Pacman emerges victorious! Score: 1357
Pacman emerges victorious! Score: 1158
Pacman emerges victorious! Score: 1338
Pacman emerges victorious! Score: 1340
Pacman emerges victorious! Score: 1161
Pacman emerges victorious! Score: 1171
Pacman emerges victorious! Score: 1252
Pacman emerges victorious! Score: 1367
Average Score: 1244.5
Scores: 1154.0, 1147.0, 1357.0, 1158.0, 1338.0, 1340.0, 1161.0, 1171.0, 1252.0, 1367.0
Win Rate: 10/10 (1.00)
Record: Win, Win, Win, Win, Win, Win, Win, Win, Win, Win
*** PASS: test_cases\part4\grade-agent.test (8 of 8 points)
*** EXTRA CREDIT: 2 points
*** 1244.5 average score (4 of 4 points)
*** Grading scheme:
*** < 500: 0 points
*** >= 500: 2 points
*** >= 1000: 4 points
*** 10 games not timed out (2 of 2 points)
*** Grading scheme:
*** < 0: fail
*** >= 0: 0 points
*** >= 5: 1 points
*** >= 10: 2 points
*** 10 wins (4 of 4 points)
*** Grading scheme:
*** < 1: fail
*** >= 1: 1 points
*** >= 4: 2 points
*** >= 7: 3 points
*** >= 10: 4 points
### Question part4: 10/10 ###

```

```

340     # Begin your code (Part 4)
341     pacman = currentGameState.getPacmanPosition()
342
343     ret = 2*currentGameState.getScore()
344     if currentGameState.isLose():
345         ret += -9999999
346     if currentGameState.isWin():
347         ret += 9999999

```

Firstly, I get the position of pacman, and get the score of that state. Then, I will check if this state is win state or loss state. If yes, I will set the return value to a very high or very low number.

```

348
349     for capsule_loc in currentGameState.getCapsules():
350         ret -= 10 * distance(pacman, capsule_loc)
351     ret += 1000 - 200 * len(currentGameState.getCapsules())

```

Secondly, in the for loop, I will let the pacman to get closer to the capsule. The 351 line will make the pacman to eat as many capsules as possible.

```

352
353     for ghost in currentGameState.getGhostStates():
354         if ghost.scaredTimer > 5:
355             ret -= 20*distance(pacman, ghost.getPosition())
356         else:
357             ret += 5*distance(pacman, ghost.getPosition())

```

In this for loop, it will check the state of ghosts, if a ghost is scared, the pacman will try to get closer to the ghost to eat it, otherwise, it will run away the ghost. Here, I set different weight since I want to let the pacman have different tendency to the ghost. In this case, the pacman will have more willing to eat the ghost than run away the ghost. In my test, I found that the pacman will allow a short distance to the ghost not scared rather run very far away. And it will crazily to get closer to the ghost get scared.

```

358
359     ret += 1000 - 20*currentGameState.getNumFood()
360
361     currentFood = currentGameState.getFood()
362     for i, d in enumerate(currentGameState.getFood()):
363         for j, e in enumerate(d):
364             if currentFood[i][j] == True:
365                 ret -= distance(pacman, (i,j))
366     return ret

```

In line 359, I will let the pacman to eat as many as food as possible. In the for loop, the pacman will try to go to the region that has more food.

Finally, return the return value.

```
367
368 def distance(l1, l2):
369     return sqrt((l1[0] - l2[0]) ** 2 + (l1[1] - l2[1]) ** 2)
370
371     # End your code (Part 4)
```

Here, I write a distance function to calculate the Euclidean distance between 2 points.

Final result:

```
Provisional grades
-----
Question part1: 25/25
Question part2: 30/30
Question part3: 30/30
Question part4: 10/10
-----
Total: 95/95
```